Implementation details of reinforcement-learning based real-time power flow management in microgrids

Rima Oulhaj, Pierre Garambois, and Lionel Roucoules

Laboratoire d'Ingénierie des Systèmes Physiques et Numériques, Arts et Métiers
Institute of Technology, HESAM Université
{rima.oulhaj, pierre.garambois, lionel.roucoules}@ensam.eu

Abstract. Real-time Power flow Management (RPM) in electric microgrids is a sequential decision making problem that can be formulated as a discounted infinite-horizon Markov Decision Process (MDP). Policy gradient methods have shown great results in solving discounted infinite-horizon MDPs in a wide range of applications. This paper provides implementation details of actor-critic policy gradient algorithms for RPM. We present the microgrid model, the MDP formulation of the RPM problem and the different test environments (i.e., instances of microgrid sizing and reward function setting labeled environment 1, 2, 3 and 4). We provide the hyperparameters of the algorithms and the algorithm optimizations used that are recommended in the literature.

1 Microgrid model

The MG is modeled as the concatenation of devices that either inject or draw power into/from the MG. At each time step, power flows inside the microgrid are balanced and verify the following equation:

$$\forall t: \quad P_t^L + P_t^{Ren} + P_{out,t}^G + P_t^B + P_t^U = 0 \tag{1}$$

Where P_t^L is the power load which verifies $P_t^L \leq 0$. P_t^{Ren} is the renewable generation given by $P_t^{Ren} = P_t^{PV} + P_t^{WT}$ where P_t^{PV} and P_t^{WT} are the photovoltaic (PV) panel and wind turbine (WT) generation respectively with $P_t^{PV} \geq 0$ and $P_t^{WT} \geq 0$. $P_{out,t}^G$ is the total generation of a set G of controllable generators given by $P_{out,t}^G = \sum_{g \in G} P_{out,t}^g$ where $\forall g, P_{out,t}^g \geq 0$. P_t^B is the equivalent power flow of a set B of batteries given by $P_t^B = \sum_{b \in B} P_t^b$ where $\forall b \in B, P_t^b \leq 0$ in the charge mode and $P_t^b \geq 0$ in the discharge mode. P_t^U is the utility grid power flow that is responsible of balancing the power flows inside the microgrid.

1.1 Controllable generation model

The power output $P_{out,t}^g$ of each controllable generator $g \in G$ at time step t is given by:

$$P_{out,t}^g = \min\left(P_N^g \cdot \mu_g, (\alpha_g \cdot \Delta t + P_{out,t-1}^g) \cdot \mu_g, P_t^g\right) \tag{2}$$

Where P_t^g is the power generation command signal (the demanded power), P_N^g is the nominal power output, μ_g is the efficiency coefficient, α_g is given by $\alpha_g = \dfrac{P_N^g}{\Delta T_g}$ and ΔT_g is the response time of $g \in G$ in $[h]$.

1.2 Battery model

For each battery $b \in B$, let W_t^b be the stored energy in $[Wh]$ at time step t:

$$W_t^b = W_{t-1}^b \cdot (1 - \sigma_{SD}^b) + (P_{C,t}^b \cdot \eta_C^b - \frac{P_{D,t}^b}{\eta_D^b}) \cdot \Delta_t \tag{3}$$

Where Δ_t is the time step length in $[h]$, W_{t-1}^b is the stored energy from the previous time step, σ_{SD}^b is the self-discharge rate, η_C^b and η_D^b are the charge and discharge efficiencies respectively, $P_{C,t}^b$ and $P_{D,t}^b$ are the charge and discharge power respectively given by $P_{C,t}^b = -\min(0, P_t^b)$, $P_{C,t}^b \geq 0$ and $P_{D,t}^b = \max(0, P_t^b)$, $P_{D,t}^b \geq 0$. For each $b \in B$ we define the state of charge (SOC) at each time step t as:

$$SOC_t^b = \frac{W_t^b}{W_{NC}^b} \tag{4}$$

Where W_{NC}^b is the nominal capacity (NC) of $b \in B$ in $[Wh]$. The SOC must verify the following constraint at each time step:

$$SOC_{min}^b \leq SOC_t^b \leq SOC_{max}^b \tag{5}$$

Where SOC_{min}^b and SOC_{max}^b are the minimum and maximum SOC respectively and are given characteristics of battery b.

1.3 Wind Turbine Model

We use a parametric model for wind turbine power curves that incorporates wind speed data [1]. Wind power production P_t^{WT} can be calculated as a function of the wind speed $V_{W,t}$, air density ρ, rotor area A_{rotor} and power coefficient $C_p(\lambda_t, \beta)$ with λ_t being the tip-speed ratio and β the blade pitch angle.

$$P_t^{WT} = \frac{1}{2}\rho \cdot A_{rotor} \cdot (V_{W,t})^3 \cdot C_p(\lambda_t, \beta) \tag{6}$$

The power coefficient $C_p(\lambda_t, \beta)$ is given by:

$$C_p(\lambda_t, \beta) = c_1(\frac{c_2}{\lambda_i} - c_3\beta - c_4\lambda_i\beta - c_5\beta^\chi - c_6)e^{-c_7/\lambda_i} + c_8\lambda_t \tag{7}$$

Where:

$$\lambda_i^{-1} = (\lambda_t + c_9\beta)^{-1} - c_{10}(\beta^3 + 1)^{-1} \tag{8}$$

Parameter	β	c_1	c_2	c_3	c_4	c_5	χ	c_6	c_7	c_8	c_9	c_{10}
Value	0	0.22	120	0.4	0	0	0	5	12.5	0	0.08	0.035

Table 1. Parameter setting for power coefficient C_p [1]

The tip-speed ratio λ_t can be formulated as a function of the rotor rotational speed and radius $D_{rotor}/2$ as well as the wind speed:

$$\lambda_t = \frac{\omega_t \cdot (D_{rotor}/2)}{V_{W,t}} \tag{9}$$

Where:

$$\omega_t = \min\left(\omega_{max}, \max(\omega_{min}, \frac{\lambda_{opt}}{D_{rotor}/2} \cdot V_{W,t})\right) \tag{10}$$

And:

$$\lambda_{opt} = \text{argmax}_{\lambda, \beta=0}\, C_p(\lambda, \beta) \tag{11}$$

Where ω_{min} and ω_{max} are the minimum and maximum rotor rotational speed respectively. We use the parameters in table 1 for equations 7 and 8. Using this model, the power curve is scaled by the nominal power of the turbine, and then the cut-in and cut-off wind speeds are applied [1].

1.4   PV panel model

The power output P_t^{PV} of the PV module is given by:

$$P_t^{PV} = A^{PV} * \mu^{PV} * I_r(t) \tag{12}$$

Where A^{PV} is the surface area of the PV module, μ^{PV} is the efficiency coefficient and $I_r(t)$ is the global horizontal solar irradiance during time step t.

1.5   Controllable power flows constraints

We define the controllable power flows constraints as the minimum and maximum bounds of the demanded generation and the storage power flows. These bounds can be written as:

$$\forall b, \forall t : \ P_{min,t}^b \leq P_t^b \leq P_{max,t}^b$$

$$\forall g, \forall t : \ 0 \leq P_t^g \leq P_{max,t}^g$$

Eq. 2 yields:

$$P_{max,t}^g = \min\left(P_N^g \cdot \mu_g, (\alpha_g \cdot \Delta t + P_{out,t-1}^g) \cdot \mu_g\right) \tag{13}$$

Eqs. 3-4-5 yield:

$$P^b_{max,t} = -\max(P^b_{C_{min},t}, -P^b_{D_{max},t}) \tag{14}$$

And:

$$P^b_{min,t} = -P^b_{C_{max},t} \tag{15}$$

$P^b_{D_{max},t}$ (eq. 16) is the maximum possible discharge power at t which must verify the minimum SOC constraint, with $P^b_{D_{max},t} \geq 0$. $P^b_{C_{min},t}$ (eq. 17) is the minimum charge power needed to maintain SOC^b_t above SOC^b_{min} at all times, i.e., to compensate for the self-discharge when necessary, in which case $P^b_{C_{min},t} \geq 0$, otherwise it is set to $-\infty$. $P^b_{C_{max},t}$ (eq. 19) is the maximum possible charge power of battery $b \in B$ at t which must verify the maximum SOC constraint, with $P^b_{C_{max},t} \geq 0$.

$$P^b_{D_{max},t} = \max\left(0, \frac{W^b_{t-1} \cdot (1 - \sigma^b_{SD}) - SOC^b_{min} \cdot W^b_{NC}}{\Delta t} \cdot \eta^s_D\right) \tag{16}$$

$$P^b_{C_{min},t} = \begin{cases} P^b_{SOC_{min},t} & \text{if} \quad W^b_{t-1} \cdot (1 - \sigma^b_{SD}) < SOC^b_{min} \cdot W^b_{NC} \\ -\infty & \text{otherwise} \end{cases} \tag{17}$$

$P^b_{SOC_{min},t}$ is the charge power necessary to maintain the energy level above $(SOC^b_{min} \cdot W^b_{NC})$ when the self-discharge alone is sufficient to violate this constraint. We calculate it as:

$$P^b_{SOC_{min},t} = \frac{SOC^b_{min} \cdot W^b_{NC} - W^b_{t-1} \cdot (1 - \sigma^b_{SD})}{\Delta t \cdot \eta^b_C} \tag{18}$$

$$P^b_{C_{max},t} = \max\left(0, \frac{SOC^b_{max} \cdot W^b_{NC} - W^b_{t-1} \cdot (1 - \sigma^b_{SD})}{\Delta t \cdot \eta^b_C}\right) \tag{19}$$

2   Power loads, meteorological and energy market data

We use real power loads data, meteorological data (global horizontal solar irradiance and wind speed) and energy market data to simulate the real-time power flow management using the infinite-time horizon MDP formulation. This means that values of power loads P^L_t (eq. 1), global horizontal solar irradiance $I_r(t)$ (eq. 12), wind speed $V_{W,t}$ (eq. 6), grid prices and feed-in tariffs C^L_t and C^F_t (eq. 26) are coming from our data collection.

The power loads data are provided by the Ecole Nationale Supérieure d'Arts et Métiers (Aix-en-Provence campus). Meteorological data are extracted using the ODRE [1] open-source API. Load costs data are extracted from the Bourses de

---

[1] https://opendata.reseaux-energies.fr/

l'électricité Spot France website[2]. The feed-in tariffs are calculated by multiplying the load costs by a coefficient (set to 0.2). Using this data and the microgrid model in section 1, the obtained PV and WT generation profiles are shown in figures 1 and 2 respectively. The power load profiles, feed-in tariffs and load costs data are shown in figures 3, 4 and 5 respectively.


Fig. 1. PV panel generation data


Fig. 2. WT generation data

**Fig. 3.** Power loads data


Fig. 4. Feed-in tariffs data

Fig. 5. Load costs data

---

[2] https://www.services-rte.com/fr/visualisez-les-donnees-publiees-par-rte/bourses-de-lelectricite-spot-france.html

3 The infinite-horizon MDP formulation of the RPM problem

3.1 State space

The state space $\mathcal{S}$ is continuous and given by:

$$\forall t \; [s_t \in \mathcal{S}] \Leftrightarrow s_t = [P_t^{Ren}, P_t^L, s_t^G, s_t^B, C_t^L, C_t^F] \tag{20}$$

Where $s_t^G = (P_{max,t}^g)_{g \in G}$ and $s_t^B = (P_{min,t}^b, P_{max,t}^b)_{b \in B}$. C_t^L and C_t^F are the load cost and the feed-in tariffs in $[\text{\euro}/Wh]$.

3.2 Action space

The action space is continuous and given by:

$$\mathcal{A} = \prod_{g \in G}[0, 1] \times \prod_{b \in B}[-1, 1] \tag{21}$$

Where:

$$a_t \in \mathcal{A} \Leftrightarrow a_t = ((a_t^g)_{g \in G}, \; (a_t^b)_{b \in B}) \tag{22}$$

$$\forall t \leq T, \; \forall b \in B, \; \forall g \in G: \;\; 0 \leq a_t^g \leq 1, \; -1 \leq a_t^b \leq 1$$

And:

$$\forall t \leq T, \forall g \in G: \; P_t^g = a_t^g \cdot P_{max,t}^g \tag{23}$$

$$\forall t, \leq T, \forall b \in B: \; P_t^b = \begin{cases} a_t^b \cdot P_{max,t}^b & if \;\;\; a_t^b \geq 0 \\ -a_t^b \cdot P_{min,t}^b & otherwise \end{cases} \tag{24}$$

3.3 Different reward functions

Economic reward function The normalized economic reward function is given by:

$$\hat{R}^E(s_t, a_t) = \frac{R^E(s_t, a_t) - R_{min}^E}{R_{max}^E - R_{min}^E} \tag{25}$$

Where R^E is defined as:

$$R^E(s_t, a_t) = C_t^F \cdot \max(0, P_t^U) + C_t^L \cdot \min(0, P_t^U) \tag{26}$$

The utility grid power flow P_t^U is given by eq. 1. R_{min}^E and R_{max}^E are empirical bounds of R^E that are computed using our training data and are given by $R_{min}^E = C_{max}^L * P_{min}^U$ and $R_{max}^E = C_{max}^F * P_{max}^U$ with:

$$P_{min}^U = P_{min}^L - \sum_{b \in B} SOC_{max}^b * W_{NC}^b \tag{27}$$

$$P_{max}^U = P_{max}^{PV} + P_{max}^{WT} + \sum_{g \in G} P_N^g + \sum_{b \in B} SOC_{max}^b * W_{NC}^b \tag{28}$$

Values $C_{max}^L, C_{max}^F, P_{max}^{PV}, P_{max}^{WT}$ and P_{min}^L are the maximum load cost, maximum feed-in tariff, maximum PV generation, maximum WT generation and minimum power load (in negative values) respectively that are derived from our training data.

Rule-based reward function This reward function is used for the particular case where the microgrid consists only of renewable energy and a battery ($G = \emptyset$, $B = \{b\}$).

Let $\mathcal{S}_1 = \{s_t \in \mathcal{S} | \delta_t > 0\}$, $\mathcal{S}_2 = \{s_t \in \mathcal{S} | \delta_t < 0\}$ and $\mathcal{S}_3 = \{s_t \in \mathcal{S} | \delta_t = 0\}$ where $\delta_t = P_t^{Ren} - P_t^L$. We define $\mathcal{S}_{11}$ and $\mathcal{S}_{12}$ as: $\mathcal{S}_{11} = \{s_t \in \mathcal{S}_1 | \delta_t \geq -P_{min,t}^b\}$, $\mathcal{S}_{12} = \{s_t \in \mathcal{S}_1 | \delta_t < -P_{min,t}^b\}$.

A partition of $\mathcal{S}$ is therefore given by $\{\mathcal{S}_{11}, \mathcal{S}_{12}, \mathcal{S}_2, \mathcal{S}_3\}$.

Sets $\{\mathcal{A}_{11}, \mathcal{A}_{12}\}$, $\{\mathcal{A}_{21}, \mathcal{A}_{22}, \mathcal{A}_{23}\}$ and $\{\mathcal{A}_{31}, \mathcal{A}_{32}\}$ are partitions of $\mathcal{A}$ where:

$$\mathcal{A}_{11} = \{a_t \in \mathcal{A} | P_t^b \geq 0\}, \mathcal{A}_{12} = \{a_t \in \mathcal{A} | P_t^b < 0\} \tag{29}$$

$$\mathcal{A}_{21} = \{a_t \in \mathcal{A} | P_t^b > 0 \wedge P_t^b < -\delta_t\}, \mathcal{A}_{22} = \{a_t \in \mathcal{A} | P_t^b > 0 \wedge P_t^b \geq -\delta_t\},$$

$$\mathcal{A}_{23} = \{a_t \in \mathcal{A} | P_t^b \leq 0\} \tag{30}$$

$$\mathcal{A}_{31} = \{a_t \in \mathcal{A} | P_t^b = 0\}, \mathcal{A}_{32} = \{a_t \in \mathcal{A} | P_t^b \neq 0\} \tag{31}$$

The rule based reward is defined as:

$$
\begin{aligned}
R^{rule}(s_t, a_t) = & \sum_{\mathcal{S}_{xx} \in \{\mathcal{S}_{11}, \mathcal{S}_{12}\}} \left[\sum_{\mathcal{A}_{yy} \in \{\mathcal{A}_{11}, \mathcal{A}_{12}\}} 1_{\mathcal{S}_{xx}}(s_t) * 1_{\mathcal{A}_{yy}}(a_t) * r_{\mathcal{S}_{xx} \times \mathcal{A}_{yy}}(s_t, a_t) \right] \\
& + \sum_{\mathcal{A}_{yy} \in \{\mathcal{A}_{21}, \mathcal{A}_{22}, \mathcal{A}_{23}\}} 1_{\mathcal{S}_2}(s_t) * 1_{\mathcal{A}_{yy}(a_t)} * r_{\mathcal{S}_2 \times \mathcal{A}_{yy}}(s_t, a_t) \\
& + \sum_{\mathcal{A}_{yy} \in \{\mathcal{A}_{31}, \mathcal{A}_{32}\}} 1_{\mathcal{S}_3}(s_t) * 1_{\mathcal{A}_{yy}(a_t)} * r_{\mathcal{S}_3 \times \mathcal{A}_{yy}}(s_t, a_t)
\end{aligned}
\tag{32}
$$

$$r_{\mathcal{S}_{11} \times \mathcal{A}_{12}}(s_t, a_t) = r_{\mathcal{S}_2 \times \mathcal{A}_{21}}(s_t, a_t) = \begin{cases} -\dfrac{P_t^b}{\delta_t} & \text{if} \quad \delta_t \neq 0 \\ 0 & \text{otherwise} \end{cases} \tag{33}$$

$$r_{\mathcal{S}_{12} \times \mathcal{A}_{12}}(s_t, a_t) = \frac{P_t^b}{P_{min,t}^b} \tag{34}$$

$$r_{\mathcal{S}_2 \times \mathcal{A}_{22}}(s_t, a_t) = \frac{1}{1 + |P_t^b + \delta_t|} \tag{35}$$

$$r_{\mathcal{S}_3 \times \mathcal{A}_{31}}(s_t, a_t) = 1 \tag{36}$$

$$r_{\mathcal{S}_{11} \times \mathcal{A}_{11}}(s_t, a_t) = r_{\mathcal{S}_{12} \times \mathcal{A}_{11}}(s_t, a_t) = r_{\mathcal{S}_2 \times \mathcal{A}_{23}}(s_t, a_t) = r_{\mathcal{S}_3 \times \mathcal{A}_{32}}(s_t, a_t) = 0 \tag{37}$$

4 Test cases: reinforcement learning environments

4.1 Microgrid sizing

Microgrid A: a microgrid with controllable generators Microgrid A includes two controllable generators labeled 'ctrl_gen_1' and 'ctrl_gen_2 (table 2), two batteries labeled 'battery_1' and 'battery_2' (table 3), a PV panel (table 5) and a WT (table 6). The variation interval of the wind turbine parameters can be found in [1].

Device label	P_N^g	μ_g	ΔT_g
Ctrl_gen_1	45000 W	0.99	$1H$
Ctrl_gen_2	15000 W	0.89	$1H$

Table 2. Controllable generator parameters — Microgrid A

Device label	σ_{SD}^b	η_C^b	η_D^b	W_{NC}^b	SOC_{min}^b	SOC_{max}^b
Battery_1	0,0239	0,97	0,82	70000 Wh	0,1	0,8
Battery_2	0,0239	0,97	0,89	60000 Wh	0,1	0,93

Table 3. Battery parameters — Microgrid A

Microgrid B: the renewables-storage microgrid This microgrid only includes renewable generation (tables 5 and 6) and a battery (table 4).

σ_{SD}^b	η_C^b	η_D^b	W_{NC}^b	SOC_{min}^b	SOC_{max}^b
0,0239	0,97	0,82	90000 Wh	0,1	0,9

Table 4. Battery parameters — Microgrid B

A^{PV}	μ^{PV}
700 m^2	$\mu^{PV} = 0.18$

Table 5. PV panel parameters — Microgrids A and B

D_{rotor}	ω_{min}	ω_{max}	Nominal power	Cut-in wind speed	Cut-out wind speed
$100\ m$	aD_{rotod}^{b} $a = 1046,558$ $b = -1,0911$	cD_{rotor}^{d} $c = 705,406$ $d = -0,8349$	$20000\ W$	$3\ ms^{-1}$	$25\ ms^{-1}$

Table 6. Wind turbine parameters — Microgrids A and B

4.2 Environments

Table 7 describes a set of environments each composed of a microgrid and a choice of reward function.

Environment label	Microgrid	Reward function
Environment 1	Microgrid A	$\hat{R}_E$
Environment 2	Microgrid B	$\hat{R}_E$
Environment 3	Microgrid B	$\alpha_E \cdot \hat{R}_E + \alpha_R \cdot R^{rule}$ $\alpha_E = 0.25$ $\alpha_R = 0.75$
Environment 4	Microgrid B	R^{rule}

Table 7. Environment definition of the RL based RPM

5 Deep RL algorithms

5.1 Algorithms VPG, TRPO and PPO

Algorithms 1, 2 and 3 provide the implementation of three actor-critic policy gradient based algorithms which work by consecutive evaluation and update of the actor parameters ω based on the objective functions J^{VPG}, J^{TRPO} and J^{PPO} in equations 38, 41 and 42 respectively.

$$J^{VPG}(\omega) = \hat{\mathbb{E}}\left[\log(\pi_\omega(s,a))\hat{A}^{\pi_\omega}(s,a)\right] \tag{38}$$

$$\hat{A}^{\pi_\omega}(s_t, a_t) = \delta_t + (\gamma\lambda)\delta_{t+1} + ... + (\gamma\lambda)^{T-t+1}\delta_{T-1} \qquad (39)$$
$$\delta_t = r_t + \gamma Q^{\pi_\omega}(s_{t+1}, a_{t+1}) - Q^{\pi_\omega}(s_t, a_t)$$

$$Q^{\pi_\omega}(s, a) = \mathbb{E}[\sum_{k=0}^{\infty} \gamma^k r_{t+k} | s_t = s, a_t = a, \pi_\omega] \qquad (40)$$

Where $\gamma \in]0, 1]$ is the discount factor and $\lambda \in]0, 1]$ is the Generalized Advantage Estimation (GAE) parameter.

$$J^{TRPO}(\omega) = \hat{\mathbb{E}}\left[r_\omega(s, a)\hat{A}^{\pi_\omega}(s, a) - \beta \cdot d_\omega(s, .) \right] \qquad (41)$$

$$J^{CLIP}(\omega) = \hat{\mathbb{E}}[\min(r_\omega(s, a)\hat{A}^{\pi_\omega}(s, a), c_\omega^\epsilon(s, a)\hat{A}^{\pi_\omega}(s, a))] \qquad (42)$$

Two neural networks are used to approximate π^ω and Q^{π_ω}. Please note that since the actor and critic networks are trained in parallel, for each training episode $k \in 0, 1, 2, ..., N - 1$ we have $Q^{\pi_{\omega_k}} = Q^{\theta_k}$ in equation 40.

Algorithm 1 NPG

Require: Initial actor network parameters ω_0
Require: Initial critic network parameters θ_0
Require: Trajectory length T
Require: Number of trajectories per training episode M
Require: Number of training episodes N
Require: Hyper-parameters λ and γ

    for $k \in 0, 1, 2, ..., N - 1$ **do**

        Sample M trajectories $(s_0^m, a_0^m, r_0^m, s_1^m, a_1^m, r_1^m, ..., s_{T-1}^m, a_{T-1}^m, r_{T-1}^m)_{m \in 1,...M}$ using actor network π_{ω_k}

        Evaluate Q-value function at each t and for each m using critic network Q^{θ_k} to compute values $\left(Q^{\theta_k}(s_t^m, a_t^m)_{t \in 0,...,T-1}\right)_{m \in 1,...,M}$

        Compute actor loss function

$$J^{VPG}(\omega_k) = \frac{1}{T \cdot M} \sum_{t=0}^{T-1} \sum_{m=1}^{M} \left[\log(\pi^{\omega_k}(s_t^m, a_t^m))\hat{A}^{\pi_{\omega_k}}(s_t^m, a_t^m) \right]$$

        Update actor network

$$\omega_{k+1} \leftarrow \text{argmax}_{\omega_k} \ J^{VPG}(\omega_k)$$

        Compute critic loss function

$$L^Q(\theta_k) = \frac{1}{(T-1) \cdot M} \sum_{t=0}^{T-2} \sum_{m=1}^{M} \left[Q^{\theta_k}(s_t^m, a_t^m) - r_t^m - \gamma Q^{\theta_k}(s_{t+1}^m, a_{t+1}^m) \right]^2$$

        Update critic network $\theta_{k+1} \leftarrow \text{argmin}_{\theta_k} \ L^Q(\theta_k)$
    end for

Algorithm 2 TRPO

Require: Initial actor network parameters ω_0
Require: Initial critic network parameters θ_0
Require: Trajectory length T
Require: Number of trajectories per training episode M
Require: Number of training episodes N
Require: Hyper-parameters λ, γ and β

Initialize old actor: $\omega_{old} \leftarrow \omega_0$
for $k \in 0, 1, 2, ..., N-1$ **do**
  Sample M trajectories $(s_0^m, a_0^m, r_0^m, s_1^m, a_1^m, r_1^m, ..., s_{T-1}^m, a_{T-1}^m, r_{T-1}^m)_{m \in 1,...M}$ using actor network π_{ω_k}
  Evaluate Q-value function at each t and for each m using critic network Q^{θ_k} to compute values $\left(Q^{\theta_k}(s_t^m, a_t^m)_{t \in 0,...,T-1}\right)_{m \in 1,...,M}$
  Compute actor loss function

$$J^{TRPO}(\omega_k) = \frac{1}{T \cdot M} \sum_{t=0}^{T-1} \sum_{m=1}^{M} \left[r_{\omega_k}(s_t^m, a_t^m) \hat{A}^{\pi_{\omega_k}}(s_t^m, a_t^m) - \beta \cdot d_{\omega_k}(s, .) \right]$$

  Update actor network

$$\omega_{k+1} \leftarrow \text{argmax}_{\omega_k} \; J^{TRPO}(\omega_k)$$

  Compute critic loss function

$$L^Q(\theta_k) = \frac{1}{(T-1) \cdot M} \sum_{t=0}^{T-2} \sum_{m=1}^{M} \left[Q^{\theta_k}(s_t^m, a_t^m) - r_t^m - \gamma Q^{\theta_k}(s_{t+1}^m, a_{t+1}^m) \right]^2$$

  Update critic network $\theta_{k+1} \leftarrow \text{argmin}_{\theta_k} \; L^Q(\theta_k)$
  $\omega_{old} \leftarrow \omega_{k+1}$
end for

Algorithm 3 PPO

Require: Initial actor network parameters ω_0
Require: Initial critic network parameters θ_0
Require: Trajectory length T
Require: Number of trajectories per training episode M
Require: Number of training episodes N
Require:
Require: Hyper-parameters λ, γ and ϵ
    Initialize old actor: $\omega_{old} \leftarrow \omega_0$
    for $k \in 0, 1, 2, ..., N-1$ **do**
        Sample M trajectories $(s_0^m, a_0^m, r_0^m, s_1^m, a_1^m, r_1^m, ..., s_{T-1}^m, a_{T-1}^m, r_{T-1}^m)_{m \in 1,...M}$ using actor network π_{ω_k}
        Evaluate Q-value function at each t and for each m using critic network Q^{θ_k} to compute values $\left(Q^{\theta_k}(s_t^m, a_t^m)_{t \in 0,...,T-1}\right)_{m \in 1,...,M}$
        Compute actor loss function

$$J^{CLIP}(\omega_k) = \frac{1}{T \cdot M} \sum_{t=0}^{T-1} \sum_{m=1}^{M} \left[\min(r_{\omega_k}(s,a)\hat{A}^{\pi_{\omega_k}}(s_t^m, a_t^m), c_{\omega_k}^{\epsilon}(s_t^m, a_t^m)\hat{A}^{\pi_{\omega_k}}(s_t^m, a_t^m)) \right]$$

        Update actor network

$$\omega_{k+1} \leftarrow \text{argmax}_{\omega_k} \ J^{CLIP}(\omega_k)$$

        Compute critic loss function

$$L^Q(\theta_k) = \frac{1}{(T-1) \cdot M} \sum_{t=0}^{T-2} \sum_{m=1}^{M} \left[Q^{\theta_k}(s_t^m, a_t^m) - r_t^m - \gamma Q^{\theta_k}(s_{t+1}^m, a_{t+1}^m) \right]^2$$

        Update critic network $\theta_{k+1} \leftarrow \text{argmin}_{\theta_k} \ L^Q(\theta_k)$
        $\omega_{old} \leftarrow \omega_{k+1}$
    end for

5.2   Hyper-parameters of algorithms VPG, TRPO and PPO

Table 8 provides the hyperparameters of algorithms 1, 2 and 3. We use the Adam optimizer in these algorithms to train the actor and critic. Following the recommendations of [2], we use the non-stationary Adam i.e., we set the parameters β_1 and β_2 of the Adam optimizer to the same value.

Hyperparameter	Value
Trajectory length T	24
Number of trajectories M	30
Number of training episodes N	$3 \cdot 10^4$
Discount factor γ	0.99
GAE parameter λ (eq. 39)	0.95
PPO clipping parameter ϵ (eq. 42)	0.2
TRPO β coefficient (eq. 41)	1
Adam learning rate	10^{-4}
Adam betas β_1 and β_2	0.9 and 0.9

Table 8. Hyperparameters of VPG, TRPO and PPO (algorithms 1, 2 and 3 respectively).

5.3   The actor and critic network

Figures 6 and 7 give the layout of the actor and critic networks. Since the output of the actor network for a given state s is the mean and standard deviation of $\pi^\omega(s,.)$, we use appropriate activation functions for the output layer (using sigmoid and tanh functions) in order to generate mean and standard deviation vectors that match our action space $\mathcal{A} = \prod_{g \in G}[0,1] \times \prod_{b \in B}[-1,1]$. An important implementation detail is that we use the truncated multivariate normal instead of the multivariate normal distribution to sample actions. Indeed, using the normal distribution introduces the possibility of infeasible actions which we want to omit.

6   Evaluating the RL algorithms: the dispersion across runs (DaR) metric

The DaR metric was proposed by [3]. Let N be the number of training episodes within one run of the RL algorithm and let L be the number of times that we run the training process. For each $l \in 1, ..., L$ we define the l-th learning curve $\mathcal{C}_l^N$ of the actor network as $\mathcal{C}_l^N = (\mathcal{R}_l^1, \mathcal{R}_l^2, ..., \mathcal{R}_l^N)$ where $\mathcal{R}_l^i$ is the cumulative reward
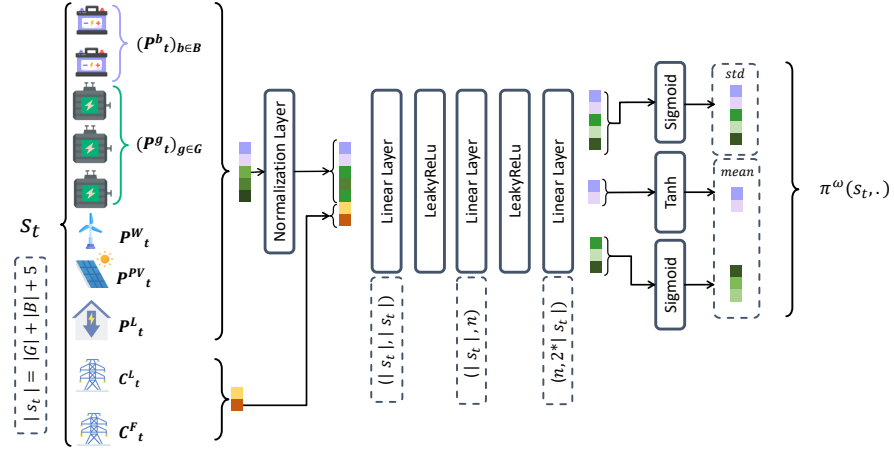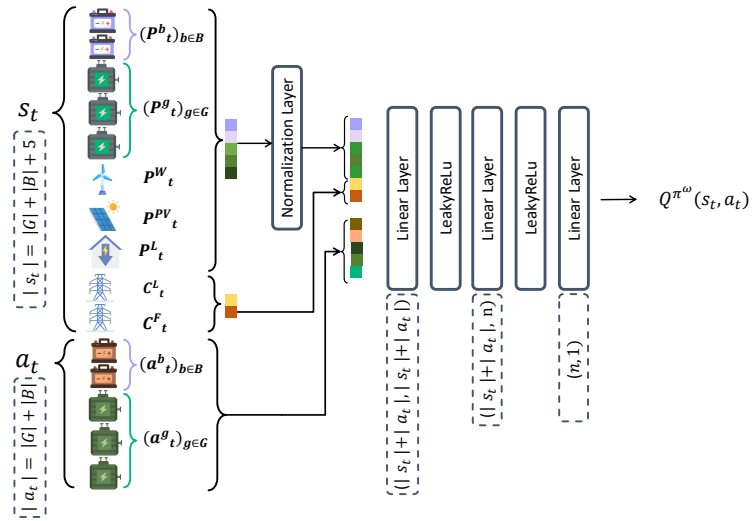
Fig. 6. Layout of the actor network

Fig. 7. Layout of the critic network

obtained after the i-th training episode of the l-th run of the training process. Low-pass filtering or smoothing of the learning curves before computing the DaR is recommended by [3]. For this purpose we use the following exponential smoothing function that is defined for a curve $\mathcal{C} = (y_1, ..., y_N)$ of length N such that $\forall i \in [1, N] : y_i \in \mathbb{R}$:

$$\text{smooth}_\varsigma(\mathcal{C}) = (\bar{y}_1, ..., \bar{y}_N) \tag{43}$$
$$\bar{y}_1 = y_1$$
$$\forall n \in 2, ..., N : \ \bar{y}_n = \varsigma \cdot y_n + (1 - \varsigma) \cdot y_{n-1}$$

Where $0 < \varsigma < 1$ is the smoothing factor. We calculate the DaR at a set of evaluation points by sampling them at a certain frequency $f \in \mathbb{N}$:

$$\text{sample}_f(\mathcal{C}) = (y_{t \cdot f})_{t \in 1, ..., \lceil \frac{N}{f} \rceil} \tag{44}$$

DaR has an across-runs variability axis, which means that it is calculated for L learning curves $(\mathcal{C}_l^N)_{l \in L}$:

$$DaR_{f,\varsigma}\left((\mathcal{C}_l^N)_{l \in L}\right) = \left(IQR(S_l^{f,\varsigma})\right)_{l \in L} \tag{45}$$
$$S_l^{f,\varsigma} = \text{sample}_f(\text{smooth}_\varsigma(\mathcal{C}_l^N))$$

7 The local optimum problem in the economic RPM of a microgrid with controllable generators

In the economic RPM problem (environment 1 in table 7), a very important observation is that the reinforcement learning algorithms reach a local optimum when the microgrid includes controllable generators. Figure 8 shows this phenomenon while using PPO. We can see that the controllable generators (labeled 'ctrl_gen_1' and 'ctrl_gen_2') are always operated near their maximum capacity and the batteries (labeled 'battery_1' and 'battery_2') are discharged in the first time step and never recharged. The controllable generators and batteries' parameters used in this experiment are given in tables 2 and 3. Note that in the test case shown, the batteries' initial energy levels are generated randomly.

Fig. 8. Economic RPM of a microgrid with controllable generators: the RL algorithm is stuck in a local optimum where the generators are operated near their maximum capacity and the batteries are never recharged.

References

[1] Yves-Marie Saint-Drenan et al. "A parametric model for wind turbine power curves incorporating environmental conditions". In: *Renewable Energy* 157 (2020), pp. 754–768. ISSN: 0960-1481. DOI: https://doi.org/10.1016/j.renene.2020.04.123.

[2] Shibhansh Dohare, Qingfeng Lan, and A. Rupam Mahmood. "Overcoming Policy Collapse in Deep Reinforcement Learning". In: *Sixteenth European Workshop on Reinforcement Learning*. 2023.

[3] Stephanie C. Y. Chan et al. *Measuring the Reliability of Reinforcement Learning Algorithms*. 2019.